

Objective: Predicting housing prices

The main objective in this notebook is to show proof of concept in accurately predicting housing prices using timeseries modeling for a real estate investment firm.

```
#importing packages and data
```

```
import pandas as pd
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import math
from math import sqrt
```

```
%matplotlib inline
```

```
import statsmodels.api as sm
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.graphics.tsaplots import plot_pacf
from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_predict
from statsmodels.tsa.arima_process import arma_generate_sample
```

```
from sklearn.metrics import mean_squared_error
```

```
df = pd.read_csv('data/zillow_data.csv.zip')
```

Data Understanding and Limitations

The data used in this project is from Zillow's housing data and is the median house price of zipcodes across the U.S.

General Overview:

- Over 14k unique zipcodes, this goes down to under 200 after we decided on which zip codes to use for our modeling
- Time range is from 1996 to April, 2018
- Prices are the median house price in the zipcode
- Data is not current and does not include anything after April 2018
- Data is very general, and we generalize even more after deciding on which zipcodes to use
- Our models are all built on the mean of the median of all the zipcode in our metros that we selected

```
#here we are filtering out our zipcodes to just look at the top 500 in
sizerank
df['Top500'] = df['SizeRank'].apply(lambda x: True if x <= 500 else
False)
```

```
df = df.loc[df.Top500, :]
```

```
#creating growth rate metric and sorting to see top zips
df['GrowthRate'] = df['2018-04'] / df['1996-04'] - 1
df.sort_values('GrowthRate', ascending = False)[:5]
```

	RegionID	RegionName	City	State	Metro	\
117	62022	11211	New York	NY	New York	
475	62027	11216	New York	NY	New York	
191	60639	7302	Jersey City	NJ	New York	
106	62026	11215	New York	NY	New York	
258	66125	20001	Washington	DC	Washington	

	CountyName	SizeRank	1996-04	1996-05	1996-06	...
\						
117	Kings	118	133200.0	132900.0	132500.0	...
475	Kings	476	146100.0	146600.0	147200.0	...
191	Hudson	192	137200.0	137800.0	138500.0	...
106	Kings	107	225700.0	227500.0	229400.0	...
258	District of Columbia	259	92000.0	92600.0	93200.0	...

	2017-09	2017-10	2017-11	2017-12	2018-01	2018-02	2018-03
2018-04 \							
117	1424700	1435300	1440500	1463100	1496100	1531100	1581900
1623700							
475	1553100	1567700	1559700	1545700	1540200	1553600	1578400
1598700							
191	1411000	1435900	1446300	1447800	1454900	1453900	1439500
1427300							
106	2244400	2266100	2275800	2287100	2288900	2265300	2244900
2243900							
258	771200	773300	777600	780500	781600	785500	791400
793300							

	Top500	GrowthRate
117	True	11.189940
475	True	9.942505
191	True	9.403061
106	True	8.941958

```
258     True     7.622826
```

```
[5 rows x 274 columns]
```

Metro Selection

So below we have the metros with the most zipcodes in our Top 500 SizeRank

```
#checking which metros have the most top500 sizerank zipcodes  
#we are not using houston and atlanta as they have lower growth rates  
than then rest
```

```
df.value_counts('Metro')[:10]
```

```
Metro  
New York                                46  
Los Angeles-Long Beach-Anaheim        32  
Chicago                                27  
Atlanta                                23  
Miami-Fort Lauderdale                 20  
Houston                                20  
Dallas-Fort Worth                     19  
Washington                            14  
San Francisco                         14  
Charlotte                             11  
dtype: int64
```

```
#showing that almost half of zip codes in our top 500 are represented  
in the top 10 metros
```

```
df.value_counts('Metro')[:10].sum()
```

```
226
```

```
#making dfs for our Metros and checking mean growthrate  
#going to use NYC, Chi, LA, MIA, and DFW highest growth rates with  
atleast 20 zipcodes
```

```
df_nyc = df[df['Metro'] == 'New York']  
df_chi = df[df['Metro'] == 'Chicago']  
df_la = df[df['Metro'] == 'Los Angeles-Long Beach-Anaheim']  
df_dfw = df[df['Metro'] == 'Dallas-Fort Worth']  
df_mia = df[df['Metro'] == 'Miami-Fort Lauderdale']
```

Growth Rate

So we are going to select our metros that have the best combination of growth rate and zipcodes in the top 500. The reason we are not using Houston and Atlanta is they both have lower growth rates than DFW and have comprable zipcode counts.

The bottom 3 metros of D.C. San Francisco, and Charlotte all had better growth rates than DFW and CHI, but we felt they did not have enough zipcodes in the top 500.

```
#showing growth rate in descending order
```

```
print('NYC:')
print(df_nyc['GrowthRate'].mean())
print('LA:')
print(df_la['GrowthRate'].mean())
print('MIA:')
print(df_mia['GrowthRate'].mean())
print('CHI:')
print(df_chi['GrowthRate'].mean())
print('DFW:')
print(df_dfw['GrowthRate'].mean())
```

```
NYC:
3.7712542878046262
LA:
3.2231950761389916
MIA:
1.9554887689343101
CHI:
1.3872262487454181
DFW:
0.9493054726931407
```

Functions

Below we have some functions we are going to use

- The first one is for converting our dfs to datetime
- the second is for melting our zipcodes creating a mean of the median of the zipcodes in our dfs
- the third one is a stationarity/dickey-fuller test check before we begin modeling

```
# function for datetimes
```

```
def get_datetimes(df):
```

```
    Takes a dataframe:
    returns only those column names that can be converted into
    datetime objects
    as datetime objects.
```

```
    NOTE number of returned columns may not match total number of
    columns in passed dataframe
```

```
    return pd.to_datetime(df.columns.values[7:272], format='%Y-%m')
```

```
#function for melting dataframes
```

```

def melt_data(df):
    """
    Takes the zillow_data dataset in wide form or a subset of the
    zillow_dataset.
    Returns a long-form datetime dataframe
    with the datetime column names as the index and the values as the
    'values' column.

    If more than one row is passes in the wide-form dataset, the
    values column
    will be the mean of the values from the datetime columns in all of
    the rows.
    """

    melted = pd.melt(df, id_vars=['RegionName', 'RegionID',
    'SizeRank', 'City', 'State', 'Metro', 'CountyName', 'Top500',
    'GrowthRate'], var_name='time')
    melted['time'] = pd.to_datetime(melted['time'],
    infer_datetime_format=True)
    melted = melted.dropna(subset=['value'])
    return melted.groupby('time').aggregate({'value': 'mean'})

#function for checking stationarity
def stationarity_check(df):

    # Calculate rolling statistics
    roll_mean = df.rolling(window=8, center=False).mean()
    roll_std = df.rolling(window=8, center=False).std()

    # Perform the Dickey Fuller Test
    dfctest = adfuller(df['value'])

    # Plot rolling statistics:
    fig = plt.figure(figsize=(12,6))
    plt.plot(df, color='blue', label='Original')
    plt.plot(roll_mean, color='red', label='Rolling Mean')
    plt.plot(roll_std, color='black', label = 'Rolling Std')
    plt.legend(loc='best')
    plt.title('Rolling Mean & Standard Deviation')
    plt.show(block=False)

    # Print Dickey-Fuller test results
    print('Results of Dickey-Fuller Test: \n')

    dfoutput = pd.Series(dfctest[0:4], index=['Test Statistic', 'p-
    value',
    '#Lags Used', 'Number of

```

```

Observations Used'])
    for key,value in dfctest[4].items():
        dfoutput['Critical Value (%)'%key] = value
    print(dfoutput)

    return None

```

Converting to datetime

#datetime in dataframes

```

get_datetimes(df_nyc)
get_datetimes(df_la)
get_datetimes(df_mia)
get_datetimes(df_dfw)
get_datetimes(df_chi)

```

```

DatetimeIndex(['1996-04-01', '1996-05-01', '1996-06-01', '1996-07-01',
               '1996-08-01', '1996-09-01', '1996-10-01', '1996-11-01',
               '1996-12-01', '1997-01-01',
               ...,
               '2017-07-01', '2017-08-01', '2017-09-01', '2017-10-01',
               '2017-11-01', '2017-12-01', '2018-01-01', '2018-02-01',
               '2018-03-01', '2018-04-01'],
              dtype='datetime64[ns]', length=265, freq=None)

```

Melting our dataframes

#melting dataframes

```

df_nyc = melt_data(df_nyc)
df_la = melt_data(df_la)
df_dfw = melt_data(df_dfw)
df_mia = melt_data(df_mia)
df_chi = melt_data(df_chi)

```

Showing the price growth of our dataframes

Here is a plot showing all our Metros mean prices from 1996 to 2018

```

plt.figure(figsize=(16, 8), dpi=150)

```

```

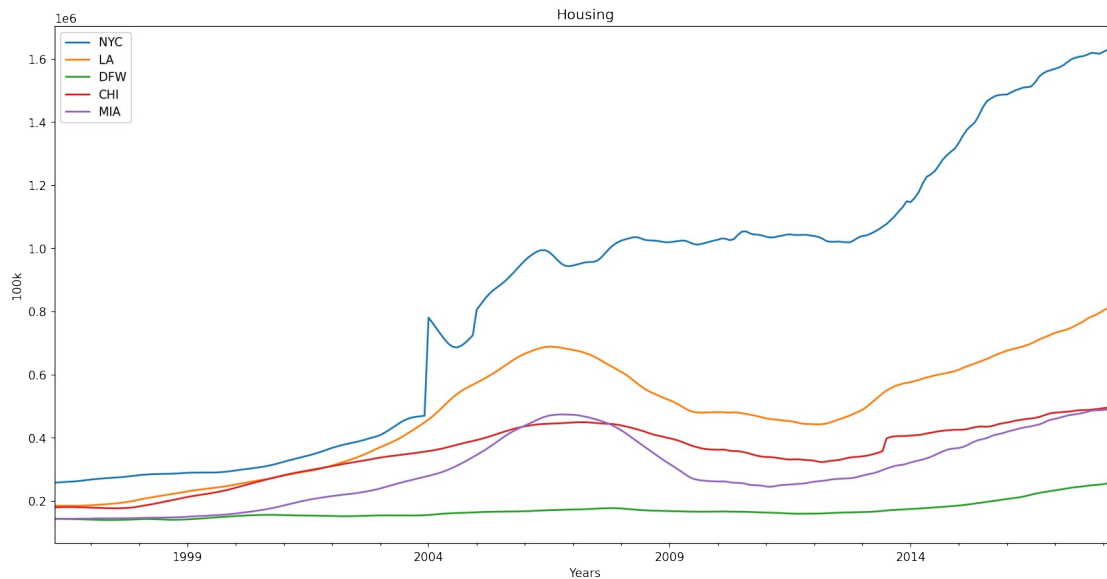
df_nyc['value'].plot(label='NYC')
df_la['value'].plot(label='LA')
df_dfw['value'].plot(label='DFW')
df_chi['value'].plot(label='CHI')
df_mia['value'].plot(label='MIA')
# adding title to the plot
plt.title('Housing')

```

adding Label to the x-axis

```
plt.xlabel('Years')
plt.ylabel('100k')
```

```
# adding legend to the curve
plt.legend();
```

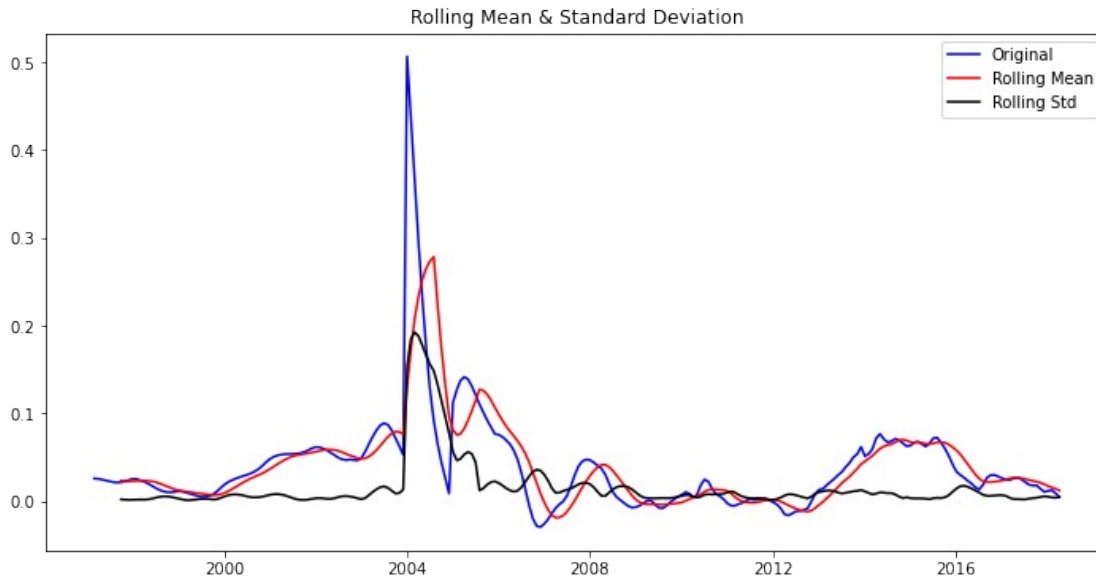


NYC ARIMA MODEL

So here is our NYC Model stationarity check, as you can see our pvalue is well below our .05 threshold so we will move forward to modeling this one.

A secondary metric here is our test statistic which we want below the critical values, this one is below all the critical values which is exactly what we want to see.

```
# Making data stationary
roll_mean_nyc = np.log(df_nyc).rolling(window=12).mean()
data_minus_roll_mean_nyc = np.log(df_nyc) - roll_mean_nyc
data_minus_roll_mean_nyc.dropna(inplace=True)
stationarity_check(data_minus_roll_mean_nyc)
```



Results of Dickey-Fuller Test:

Test Statistic	-4.358815
p-value	0.000351
#Lags Used	0.000000
Number of Observations Used	253.000000
Critical Value (1%)	-3.456464
Critical Value (5%)	-2.873033
Critical Value (10%)	-2.572895
dtype:	float64

NYC ARIMA

Below we have our train test split, our arima model's p,d,q orders, as well as our predictions for visualizations farther down.

We mainly focused on keeping our models AR and MA coef more than .2 away from zero and the pvalues under .05, as you can see this one meets those limitations.

Some of our secondary focuses were on Heteroskedasticity, Skew, and Kurtosis.

#Model DF, train/test, storing model pred for visuals, and model summary

```
model_nyc = data_minus_roll_mean_nyc
train = model_nyc[:-13]
test = model_nyc[-13:]
```

```
arima_nyc = ARIMA(train, order=(3,3,1))
arima_nyc_fit = arima_nyc.fit()
pred_nyc = arima_nyc_fit.predict(start="2017-04-01", end="2018-04-01")
pred = arima_nyc_fit.get_prediction(start="2015-04-01", end="2018-04-
```



```
01", dynamic=False)
pred_ci = pred.conf_int()
print(arima_nyc_fit.summary())
```

```
C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\
tsa\base\tsa_model.py:524: ValueWarning: No frequency information was
provided, so inferred frequency MS will be used.
```

```
warnings.warn('No frequency information was'
C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\
tsa\base\tsa_model.py:524: ValueWarning: No frequency information was
provided, so inferred frequency MS will be used.
```

```
warnings.warn('No frequency information was'
C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\
tsa\base\tsa_model.py:524: ValueWarning: No frequency information was
provided, so inferred frequency MS will be used.
```

```
warnings.warn('No frequency information was'
```

SARIMAX Results

```
=====
=====
```

```
Dep. Variable:          value    No. Observations:
241
Model:                ARIMA(3, 3, 1)    Log Likelihood
440.799
Date:                 Mon, 31 Oct 2022    AIC
-871.598
Time:                 17:50:38    BIC
-854.236
Sample:              03-01-1997    HQIC
-864.601
                   - 03-01-2017
```

```
Covariance Type:          opg
```

```
=====
=====
```

	coef	std err	z	P> z	[0.025
0.975]					

ar.L1	-0.7241	0.021	-34.466	0.000	-0.765
-0.683					
ar.L2	-0.4675	0.029	-16.269	0.000	-0.524
-0.411					
ar.L3	-0.2337	0.025	-9.233	0.000	-0.283
-0.184					
ma.L1	-0.9984	0.300	-3.331	0.001	-1.586
-0.411					
sigma2	0.0014	0.000	3.569	0.000	0.001

0.002

```
=====
=====
Ljung-Box (L1) (Q):                0.55    Jarque-Bera (JB):
101130.92
Prob(Q):                0.46    Prob(JB):
0.00
Heteroskedasticity (H):            3.32    Skew:
6.73
Prob(H) (two-sided):            0.00    Kurtosis:
103.09
=====
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
warnings.warn("Maximum Likelihood optimization failed to "

Visualization showing predications, confidence interval, and actual values


#Predictions, CI, and actual values visual

```
ax = test.plot(label='Actual values')
pred.predicted_mean.plot(ax=ax, label='Predictions', alpha=.7,
color='red')
```

```
ax.fill_between(pred_ci.index,
                pred_ci.iloc[:, 0],
                pred_ci.iloc[:, 1], color='k', alpha=.2)
```

```
ax.set_xlabel('Date')
ax.set_ylabel('Price')
plt.legend()
```

```
plt.show()
```




Zoomed in predictions and actual values

#zoomed in pred and value

```
plt.plot(test)
```

```
plt.plot(pred_nyc, color= 'red')
```

```
[<matplotlib.lines.Line2D at 0x23fb6f612e0>]
```



MSE, RMSE, and Residual check

All of these are similar metrics and once again we are trying to get them as close to zero as possible

```
#MSE and RMSE check
```

```
expected = test
predictions = pred_nyc
mse = mean_squared_error(expected, predictions)
rmse = sqrt(mse)
print('MSE: %f' % mse)
print('RMSE: %f' % rmse)
```

```
MSE: 0.000008
```


```
RMSE: 0.002805
```


```
#residual check
```

```
residuals = pd.DataFrame(arima_nyc_fit.resid)
residuals.plot()
plt.show()
```

```
residuals.plot(kind='kde')
plt.show()
```

```
print(residuals.describe())
```





```

count    241.000000
mean     -0.000015
std       0.037494
min      -0.197097
25%      -0.002194
50%       0.000232
75%       0.002258
max       0.462894

```

ACF and PACF


These ACF and PACF are for helping us tune our models p and q, while we might have not used the exact lags on this chart, it provided a starting point to tune our models and work from.

```

#ACF and PACF for model tuning
fig, ax = plt.subplots(figsize=(16,3))
plot_acf(model_nyc, ax=ax, lags=40);

fig, ax = plt.subplots(figsize=(16,3))
plot_pacf(model_nyc, ax=ax, lags=40);

```




LA Model

So here is our LA Model stationarity check, as you can see our pvalue is below our .05 threshold.

A secondary metric here is our test statistic which we want below the critical values, while this one is not below the Critical Value of 1% it is below the other two so we will move forward.

#Model stationarity

```
exp_roll_mean_la = np.log(df_la).ewm(alpha = .005).mean()
data_minus_exp_roll_mean_la = np.log(df_la) - exp_roll_mean_la
stationarity_check(data_minus_exp_roll_mean_la)
```



Results of Dickey-Fuller Test:

```
Test Statistic      -2.863459
p-value             0.049765
#Lags Used          15.000000
Number of Observations Used  249.000000
Critical Value (1%)  -3.456888
Critical Value (5%)  -2.873219
Critical Value (10%) -2.572994
dtype: float64
```

LA ARIMA

Below we have our train test split, our arima model's p,d,q orders, as well as our predictions for visualizations farther down.

We mainly focused on keeping our models AR and MA coef more than .2 away from zero and the pvalues under .05, as you can see this one meets those limitations.

Some of our secondary focuses were on Heteroskedasticity, Skew, and Kurtosis.

```
#Model DF, train/test, storing model pred for visuals, and model summary
model_la = data_minus_exp_roll_mean_la
train = model_la[:-13]
test = model_la[-13:]

arima_la = ARIMA(train, order=(1,2,1))
arima_la_fit = arima_la.fit()
pred_la = arima_la_fit.predict(start="2017-04-01", end="2018-04-01")
pred = arima_la_fit.get_prediction(start="2015-04-01", end="2018-04-01", dynamic=False)
pred_ci = pred.conf_int()
print(arima_la_fit.summary())
```

SARIMAX Results

```
=====
=====
Dep. Variable:          value    No. Observations:
252
Model:                  ARIMA(1, 2, 1)    Log Likelihood
1359.797
Date:                  Mon, 31 Oct 2022    AIC
2713.595
Time:                  17:50:39    BIC
2703.030
Sample:                04-01-1996    HQIC
2709.343
```

- 03-01-2017

Covariance Type: opg

=====					
=====					
	coef	std err	z	P> z	[0.025
0.975]					

ar.L1	0.2927	0.065	4.534	0.000	0.166
0.419					
ma.L1	0.4811	0.056	8.610	0.000	0.372
0.591					
sigma2	1.1e-06	7.01e-08	15.697	0.000	9.63e-07
1.24e-06					
=====					
=====					
Ljung-Box (L1) (Q):			0.00	Jarque-Bera (JB):	
77.13					
Prob(Q):			0.96	Prob(JB):	
0.00					
Heteroskedasticity (H):			7.28	Skew:	
0.41					
Prob(H) (two-sided):			0.00	Kurtosis:	
5.60					
=====					
=====					

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\tsa\base\tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

warnings.warn('No frequency information was')

C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\tsa\base\tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

warnings.warn('No frequency information was')

C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\tsa\base\tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

warnings.warn('No frequency information was')

Visualization showing predications, confidence interval, and actual values


#Prediction, true values, and CI graph

```
ax = test.plot(label='Actual values')
pred.predicted_mean.plot(ax=ax, label='Predictions', alpha=.7,
color='red')
```

```
ax.fill_between(pred_ci.index,
                pred_ci.iloc[:, 0],
                pred_ci.iloc[:, 1], color='k', alpha=.2)
```

```
ax.set_xlabel('Date')
ax.set_ylabel('Price')
plt.legend()
```


```
plt.show()
```



Zoomed in predictions and actual values

#zoomed in predictions/values

```
plt.plot(test)
plt.plot(pred_la, color = 'red');
```



MSE, RMSE, and Residual check

All of these are similar metrics and once again we are trying to get them as close to zero as possible

#MSE and RMSE check

```
expected = test
predictions = pred_la
mse = mean_squared_error(expected, predictions)
rmse = sqrt(mse)
print('MSE: %f' % mse)
print('RMSE: %f' % rmse)
```

MSE: 0.000851

RMSE: 0.029163

#checking residuals



```
residuals = pd.DataFrame(arima_la_fit.resid)
residuals.plot()
plt.show()
```

density plot of residuals

```
residuals.plot(kind='kde')
plt.show()
```

summary stats of residuals

```
print(residuals.describe())
```



```

0
count 252.000000
mean  0.000002
std   0.001048
min   -0.003731
25%   -0.000647
50%    0.000026
75%    0.000491
max    0.004375

```


ACF and PACF

These ACF and PACF are for helping us tune our models p and q, while we might have not used the exact lags on this chart, it provided a starting point to tune our models and work from.

```
#ACF and PACF for model tuning
```

```
fig, ax = plt.subplots(figsize=(16,3))  
plot_acf(model_la, ax=ax, lags=40);
```

```
fig, ax = plt.subplots(figsize=(16,3))  
plot_pacf(model_la, ax=ax, lags=40);
```



DFW model


So here is our DFW Model stationarity check, as you can see our pvalue is right at our .05 threshold.

A secondary metric here is our test statistic which we want below the critical values, while this is not below our 1% and 10% Critical Value, it is near enough to move forward with.

```
#Stationarity
```

```
exp_roll_mean_dfw = np.log(df_dfw).ewm(halflife = .5).mean()
```

```
data_minus_exp_roll_mean_dfw = np.log(df_dfw) - exp_roll_mean_dfw  
stationarity_check(data_minus_exp_roll_mean_dfw)
```



Results of Dickey-Fuller Test:

Test Statistic	-2.829496
p-value	0.054182
#Lags Used	5.000000
Number of Observations Used	259.000000
Critical Value (1%)	-3.455853
Critical Value (5%)	-2.872765
Critical Value (10%)	-2.572752
dtype:	float64

DFW ARIMA

Below we have our train test split, our arima model's p,d,q orders, as well as our predictions for visualizations farther down.

We mainly focused on keeping our models AR and MA coef more than .2 away from zero and the pvalues under .05. The only place where it breaks those thresholds is at the MA 3 which still has a pvalue sub .05, so we will keep it in.

Some of our secondary focuses were on Heteroskedasticity, Skew, and Kurtosis.

#Model DF, train/test, storing model pred for visuals, and model summary

```
model_dfw = data_minus_exp_roll_mean_dfw
train = model_dfw.iloc[:-13]
test = model_dfw.iloc[-13:]
```

```
arima_dfw = ARIMA(train, order=(1,1,3))
arima_dfw_fit = arima_dfw.fit()
pred_dfw = arima_dfw_fit.predict(start="2017-04-01", end="2018-04-01")
pred = arima_dfw_fit.get_prediction(start="2015-04-01", end="2018-04-01")
```

```
01", dynamic=False)
pred_ci = pred.conf_int()
print(arima_dfw_fit.summary())
```

SARIMAX Results

```
=====
=====
Dep. Variable:          value    No. Observations:
252
Model:                ARIMA(1, 1, 3)    Log Likelihood
1832.230
Date:                Mon, 31 Oct 2022    AIC                -
3654.461
Time:                17:50:40    BIC                -
3636.833
Sample:                04-01-1996    HQIC               -
3647.367
                        - 03-01-2017
```

Covariance Type: opg

```
=====
=====
                        coef    std err          z      P>|z|      [0.025
0.975]
-----
-----
ar.L1          0.6634      0.034     19.649     0.000      0.597
0.730
ma.L1          0.5334      0.020     26.632     0.000      0.494
0.573
ma.L2         -0.3308      0.021    -15.388     0.000     -0.373
-0.289
ma.L3         -0.1332      0.027     -4.941     0.000     -0.186
-0.080
sigma2      2.592e-08    1.57e-09     16.518     0.000    2.28e-08
2.9e-08
```

```
=====
=====
Ljung-Box (L1) (Q):          0.08    Jarque-Bera (JB):
132.23
Prob(Q):          0.77    Prob(JB):
0.00
Heteroskedasticity (H):      3.23    Skew:
-0.37
Prob(H) (two-sided):      0.00    Kurtosis:
6.48
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 3.04e+16. Standard errors may be unstable.

C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\tsa\base\tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

warnings.warn('No frequency information was')

C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\tsa\base\tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

warnings.warn('No frequency information was')

C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\tsa\base\tsa_model.py:524: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

warnings.warn('No frequency information was')

Visualization showing predications, confidence interval, and actual values

#Predictions, True values, and CI graph

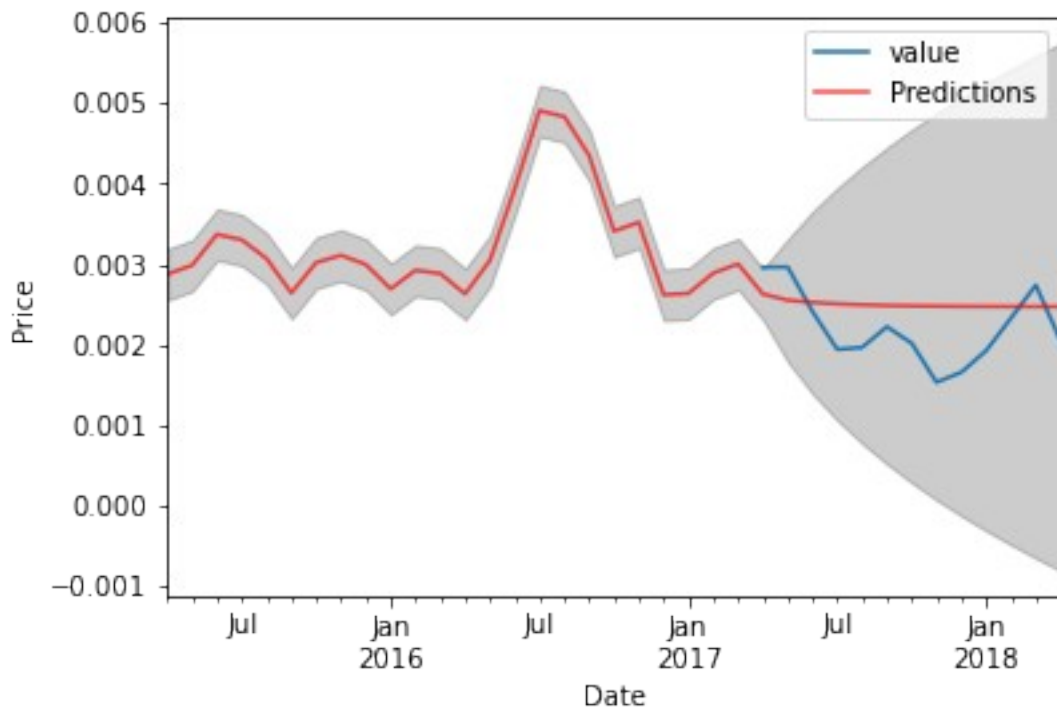
ax = test.plot(label='Actual values')

pred.predicted_mean.plot(ax=ax, label='Predictions', alpha=.7, color='red')

ax.fill_between(pred_ci.index,
 pred_ci.iloc[:, 0],
 pred_ci.iloc[:, 1], color='k', alpha=.2)

ax.set_xlabel('Date')
ax.set_ylabel('Price')
plt.legend()

plt.show()




Zoomed in predictions and actual values

#zoomed in of pred and values

```
plt.plot(test)
```

```
plt.plot(pred_dfw, color = 'red')
```

[<matplotlib.lines.Line2D at 0x23fb6f911f0>]



MSE, RMSE, and Residual check

All of these are similar metrics and once again we are trying to get them as close to zero as possible

#MSE and RMSE check

```
expected = test
predictions = pred_dfw
mse = mean_squared_error(expected, predictions)
rmse = sqrt(mse)
print('MSE: %f' % mse)
print('RMSE: %f' % rmse)
```

MSE: 0.000000

RMSE: 0.000507

#Residuals check


```
residuals = pd.DataFrame(arima_dfw_fit.resid)
residuals.plot()
plt.show()
```


density plot of residuals

```
residuals.plot(kind='kde')
plt.show()
```

summary stats of residuals

```
print(residuals.describe())
```





	0
count	252.000000
mean	0.000003
std	0.000163
min	-0.000851
25%	-0.000081
50%	0.000005
75%	0.000089
max	0.000605


ACF and PACF

#PACF and ACF for model tuning

```
fig, ax = plt.subplots(figsize=(16,3))
plot_acf(model_dfw, ax=ax, lags=40);
```

```
fig, ax = plt.subplots(figsize=(16,3))
plot_pacf(model_dfw, ax=ax, lags=40);
```

```
C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\
regression\linear_model.py:1434: RuntimeWarning: invalid value
encountered in sqrt
  return rho, np.sqrt(sigmasq)
```



Chicago Model

So here is our Chicago Model stationarity check, as you can see our pvalue is right at our .05 threshold.

A secondary metric here is our test statistic which we want below the critical values, while this is not below our 1% it is quite close and below the other critical values, so we will move forward.

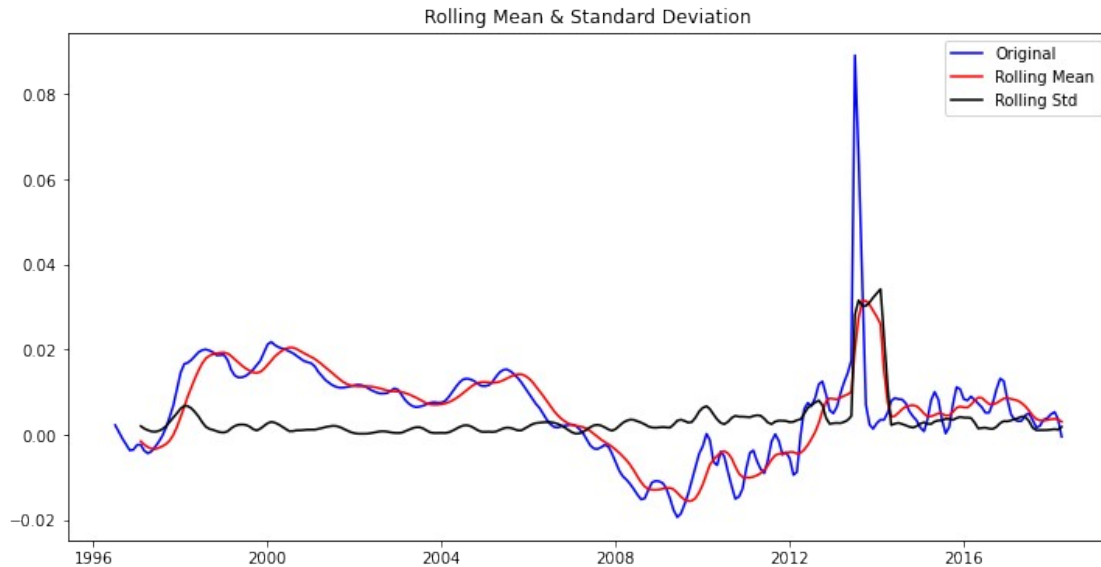
#Stationarity check

```
roll_mean_chi = np.log(df_chi).rolling(window=4).mean()
data_minus_roll_mean_chi = np.log(df_chi) - roll_mean_chi
```

Print the first 10 rows

```
data_minus_roll_mean_chi.head(10)
data_minus_roll_mean_chi.dropna(inplace=True)
```

```
stationarity_check(data_minus_roll_mean_chi)
```



Results of Dickey-Fuller Test:

Test Statistic	-3.281159
p-value	0.015742
#Lags Used	3.000000
Number of Observations Used	258.000000
Critical Value (1%)	-3.455953
Critical Value (5%)	-2.872809
Critical Value (10%)	-2.572775
dtype:	float64

Chicago ARIMA

Below we have our train test split, our arima's p,d,q orders, as well as our predictions for visualizations farther down.

We mainly focused on keeping our models AR and MA coef more than .2 away from zero and the pvalues under .05. The place's where it breaks those thresholds is at the MA 1 which has a pvalue of .12 and AR 4 which has a pvalue of .07, we are keeping it in because the other MA and AR's are sub .05

Some of our secondary focuses were on Heteroskedasticity, Skew, and Kurtosis.

```
#Model DF, train/test, storing model pred for visuals, and model
summary
model_chi = data_minus_roll_mean_chi
train = model_chi[:-13]
test = model_chi[-13:]

arima_chi = ARIMA(train, order=(4,2,3))
arima_chi_fit = arima_chi.fit()
```

```

pred_chi = arima_chi_fit.predict(start="2017-04-01", end="2018-04-01")
pred = arima_chi_fit.get_prediction(start="2015-04-01", end="2018-04-01", dynamic=False)
pred_ci = pred.conf_int()

```

```

print(arima_chi_fit.summary())

```

```

C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\
tsa\base\tsa_model.py:524: ValueWarning: No frequency information was
provided, so inferred frequency MS will be used.

```

```

warnings.warn('No frequency information was'

```

```

C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\
tsa\base\tsa_model.py:524: ValueWarning: No frequency information was
provided, so inferred frequency MS will be used.

```

```

warnings.warn('No frequency information was'

```

```

C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\
tsa\base\tsa_model.py:524: ValueWarning: No frequency information was
provided, so inferred frequency MS will be used.

```

```

warnings.warn('No frequency information was'

```

SARIMAX Results

```

=====
=====
Dep. Variable:          value    No. Observations:
249
Model:                ARIMA(4, 2, 3)    Log Likelihood
925.073
Date:                Mon, 31 Oct 2022    AIC                -
1834.146
Time:                17:50:41    BIC                -
1806.071
Sample:                07-01-1996    HQIC               -
1822.843
                    - 03-01-2017

```

```

Covariance Type:          opg

```

```

=====
=====

```

	coef	std err	z	P> z	[0.025
0.975]					

ar.L1	-1.4069	0.356	-3.949	0.000	-2.105
-0.709					
ar.L2	-0.7447	0.264	-2.824	0.005	-1.262
-0.228					
ar.L3	-0.2405	0.121	-1.992	0.046	-0.477
-0.004					

ar.L4	-0.1948	0.109	-1.782	0.075	-0.409
0.020					
ma.L1	0.5298	0.356	1.489	0.136	-0.167
1.227					
ma.L2	-0.7148	0.295	-2.422	0.015	-1.293
-0.136					
ma.L3	-0.7425	0.314	-2.368	0.018	-1.357
-0.128					
sigma2	3.221e-05	4.49e-07	71.783	0.000	3.13e-05
3.31e-05					

```

=====
=====
Ljung-Box (L1) (Q):                0.38   Jarque-Bera (JB):
113354.27
Prob(Q):                0.54   Prob(JB):
0.00
Heteroskedasticity (H):        70.82   Skew:
6.87
Prob(H) (two-sided):          0.00   Kurtosis:
107.05
=====
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals

warnings.warn("Maximum Likelihood optimization failed to "

Visualization showing predications, confidence interval, and actual values


#Prediction, true values, and CI graph

```
ax = test.plot(label='Actual values')
pred.predicted_mean.plot(ax=ax, label='Predictions', alpha=.7,
color='red')
```

```
ax.fill_between(pred_ci.index,
                pred_ci.iloc[:, 0],
                pred_ci.iloc[:, 1], color='k', alpha=.2)
```

```
ax.set_xlabel('Date')
ax.set_ylabel('Price')
plt.legend()
```

```
plt.show()
```




Zoomed in predictions and values

#zoomed in pred and values

```
plt.plot(test)
```

```
plt.plot(pred_chi, color = 'red')
```

```
[<matplotlib.lines.Line2D at 0x23fb7b31f40>]
```



MSE, RMSE, and residuals check

#MSE and RMSE check

```
expected = test
predictions = pred_chi
mse = mean_squared_error(expected, predictions)
rmse = sqrt(mse)
print('MSE: %f' % mse)
print('RMSE: %f' % rmse)
```

MSE: 0.000003

RMSE: 0.001743

#residuals check


```
residuals = pd.DataFrame(arima_chi_fit.resid)
residuals.plot()
plt.show()
```


density plot of residuals

```
residuals.plot(kind='kde')
plt.show()
```

summary stats of residuals

```
print(residuals.describe())
```





```

count    249.000000
mean     -0.000008
std       0.005672
min      -0.032479
25%      -0.001083
50%      -0.000161
75%       0.000802
max       0.071362

```

PACF and ACF

```


#PACF and ACF graph for model tuning
fig, ax = plt.subplots(figsize=(16,3))
plot_acf(model_chi, ax=ax, lags=40);


```

```

fig, ax = plt.subplots(figsize=(16,3))
plot_pacf(model_chi, ax=ax, lags=40);

```






Miami Model

So here is our Miami stationarity check, as you can see our pvalue is right at our .05 threshold.

A secondary metric here is our test statistic which we want below the critical values, while this is not below our 1% it is quite close and below the other critical values, so we will move forward.

#Stationarity

```
exp_roll_mean_mia = np.log(df_mia).ewm(alpha=0.05).mean()
data_minus_exp_roll_mean_mia = np.log(df_mia) - exp_roll_mean_mia
stationarity_check(data_minus_exp_roll_mean_mia)
```



Results of Dickey-Fuller Test:

Test Statistic	-3.062570
p-value	0.029450
#Lags Used	15.000000
Number of Observations Used	249.000000
Critical Value (1%)	-3.456888
Critical Value (5%)	-2.873219

Critical Value (10%) -2.572994
dtype: float64

Miami ARIMA

Below we have our train test split, our arima's p,d,q orders, as well as our predictions for visualizations farther down.

We mainly focused on keeping our models AR and MA coef more than .2 away from zero and the pvalues under .05. This one keeps all those within our acceptable range.

Some of our secondary focuses were on Heteroskedasticity, Skew, and Kurtosis.

#Model DF, train/test, storing model pred for visuals, and model summary

```
model_mia = data_minus_exp_roll_mean_mia
train = model_mia[:-13]
test = model_mia[-13:]
```

```
arima_mia = ARIMA(train, order=(2,0,1))
arima_mia_fit = arima_mia.fit()
pred_mia = arima_mia_fit.predict(start="2017-04-01", end="2018-04-01")
pred = arima_mia_fit.get_prediction(start="2015-04-01", end="2018-04-01", dynamic=False)
pred_ci = pred.conf_int()
print(arima_mia_fit.summary())
```

```
C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\
tsa\base\tsa_model.py:524: ValueWarning: No frequency information was
provided, so inferred frequency MS will be used.
```

```
warnings.warn('No frequency information was')
```

```
C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\
tsa\base\tsa_model.py:524: ValueWarning: No frequency information was
provided, so inferred frequency MS will be used.
```

```
warnings.warn('No frequency information was')
```

```
C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\
tsa\base\tsa_model.py:524: ValueWarning: No frequency information was
provided, so inferred frequency MS will be used.
```

```
warnings.warn('No frequency information was')
```

SARIMAX Results

```
=====
=====
Dep. Variable:                value    No. Observations:
252
Model:                ARIMA(2, 0, 1)    Log Likelihood
1228.862
Date:                Mon, 31 Oct 2022    AIC
2447.723
```

```

Time:                17:50:43    BIC                -
2430.076
Sample:             04-01-1996    HQIC                -
2440.622
                    - 03-01-2017

```

Covariance Type: opg

```

=====
=====
              coef      std err          z      P>|z|      [0.025
0.975]
-----
-----
const          0.0748      0.062       1.203      0.229      -0.047
0.196
ar.L1          1.9313      0.023     83.590      0.000       1.886
1.977
ar.L2         -0.9342      0.023    -40.453      0.000      -0.979
-0.889
ma.L1          0.4219      0.047       9.055      0.000       0.331
0.513
sigma2      3.251e-06    2.38e-07     13.669      0.000     2.78e-06
3.72e-06
=====
=====
Ljung-Box (L1) (Q):                14.73    Jarque-Bera (JB):
52.87
Prob(Q):                0.00    Prob(JB):
0.00
Heteroskedasticity (H):            22.95    Skew:
0.33
Prob(H) (two-sided):            0.00    Kurtosis:
5.15
=====
=====

```

Warnings:

```

[1] Covariance matrix calculated using the outer product of gradients
(complex-step).

```

Visualization showing predications, confidence interval, and actual values

#Predictions, values, and CI graph

```

ax = test.plot(label='Actual values')
pred.predicted_mean.plot(ax=ax, label='Predictions', alpha=.7,
color='red')

```


```

ax.fill_between(pred_ci.index,
                pred_ci.iloc[:, 0],
                pred_ci.iloc[:, 1], color='k', alpha=.2)

ax.set_xlabel('Date')
ax.set_ylabel('Price')
plt.legend()

plt.show()

```



Zoomed in predictions and values


#zoomed in pred and values

```

plt.plot(test)
plt.plot(pred_mia, color = 'red')

```

[<matplotlib.lines.Line2D at 0x23fb7b870d0>]



MSE, RMSE, and residual check

#MSE and RMSE check

```
expected = test
predictions = pred_mia
mse = mean_squared_error(expected, predictions)
rmse = sqrt(mse)
print('MSE: %f' % mse)
print('RMSE: %f' % rmse)
```

```
MSE: 0.000127
RMSE: 0.011270
```

#residuals check



```
residuals = pd.DataFrame(arima_mia_fit.resid)
residuals.plot()
plt.show()
```

density plot of residuals

```
residuals.plot(kind='kde')
plt.show()
```

summary stats of residuals

```
print(residuals.describe())
```



```

0
count 252.000000
mean  -0.000270
std    0.005045
min    -0.074753
25%    -0.000792
50%     0.000018
75%     0.000839
max     0.006774

```

PACF and ACF

#PACF and ACF graph for model tuning

```
fig, ax = plt.subplots(figsize=(16,3))  
plot_acf(model_mia, ax=ax, lags=40);
```

```
fig, ax = plt.subplots(figsize=(16,3))  
plot_pacf(model_mia, ax=ax, lags=40);
```

```
C:\Users\Clay\anaconda3\envs\learn-env\lib\site-packages\statsmodels\  
regression\linear_model.py:1434: RuntimeWarning: invalid value  
encountered in sqrt  
    return rho, np.sqrt(sigmatq)
```

